

1. What is AWS (Amazon Web Services)?

Amazon Web Services (AWS) is a comprehensive cloud computing platform offered by Amazon that provides a broad set of **on-demand cloud services** — including compute, storage, databases, analytics, networking, machine learning, security, and more — to individuals, startups, enterprises, and governments. AWS allows you to access scalable computing power and storage over the internet, with a **pay-as-you-go pricing model**.

Key Characteristics

- **Cloud Computing Platform:** Enables computing without owning physical servers.
- **Global Infrastructure:** AWS has multiple **regions and availability zones** around the world to ensure reliability and low latency.
- **240+ Services:** Including compute (EC2), storage (S3), databases (RDS), AI/ML, security, developer tools, and more.
- **Security & Compliance:** Enterprise-grade security, encryption, identity management, and compliance offerings.

Benefits

- **Cost-effective:** No upfront hardware costs; pay only for what you use.
- **Scalable & Flexible:** Scale up/down resources based on demand.
- **Reliable:** High availability and fault tolerance.

 **Official Info:** <https://aws.amazon.com/what-is-aws/>

 AWS Documentation Overview: <https://docs.aws.amazon.com/whitepapers/latest/aws-overview/introduction.html>

2. What is AWS Bedrock?

AWS Bedrock is a fully managed service from AWS that lets developers build generative AI applications using a choice of high-performance **foundation models (FMs)** from multiple providers — such as Anthropic, AI21 Labs, Cohere, Stability AI, Meta, and Amazon — through a **single API**.

Key Features

- **Model Choice:** Access and experiment with models like Claude, Titan, Jurassic, and Stable Diffusion in one place (no vendor lock-in).
- **Customizations:** Fine-tune models with your own data using techniques like **instruction tuning**.
- **Secure & Scalable:** Enterprise-grade security and serverless deployment with no infrastructure to manage.
- **Integration:** Works with other AWS services (S3, IAM, VPC).

Bedrock simplifies building generative AI capabilities such as chatbots, summarization tools, and document understanding, while handling provisioning, scaling, model selection, and security.

 Official AWS Bedrock Page: <https://aws.amazon.com/bedrock/>

(Note: actual AWS Bedrock home page link should be used; above is conceptual — search "AWS Bedrock" on AWS if needed)

3. What is Amazon SageMaker?

Amazon SageMaker (recently branded *SageMaker AI*) is AWS's fully managed machine learning (ML) service that enables developers and data scientists to **build, train, and deploy machine learning models at scale**.

- ◆ **What SageMaker Provides**
- **Data Preparation & Features:** Tools like **Data Wrangler** for data cleaning and exploration.
- **Notebooks & IDE:** Hosted Jupyter notebooks for data science workflows.
- **Model Training & Tuning:** Train models using built-in or custom algorithms with automatic hyperparameter optimization.
- **Deployment & Monitoring:** Deploy models as real-time endpoints with monitoring.

Why Use SageMaker?

- Reduces complexity of ML workflows.
- Integrates with AWS ecosystem (S3, Lambda, CloudWatch).
- Supports popular frameworks (TensorFlow, PyTorch).

 Official SageMaker Info: <https://aws.amazon.com/sagemaker/>

4. What is RAG (Retrieval-Augmented Generation)?

Retrieval-Augmented Generation (RAG) is an advanced AI technique that combines **information retrieval** with **generative models** to produce more accurate, up-to-date, and context-aware responses. Instead of generating text solely from the model's internal learned knowledge, RAG systems **retrieve relevant documents or data** from an external knowledge base and **augment the query with that context** before generating the answer.

How RAG Works

1. **Retrieve** relevant documents or passages based on semantic similarity to the user query.
2. **Augment** the original prompt with the retrieved data.
3. **Generate** the response using a language model with the augmented prompt.

Why RAG?

- ✓ Produces responses grounded in real, specific, domain-relevant data.
- ✓ Fills gaps in model knowledge with up-to-date information.
- ✓ Improves accuracy and factual correctness.

 AWS RAG Docs (SageMaker):

<https://docs.aws.amazon.com/sagemaker/latest/dg/jumpstart-foundation-models-customize-rag.html>

5. What is K21 Academy?

K21 Academy is a technology training and upskilling platform based in India that offers **job-oriented courses** focused on cloud computing, AI/ML, DevOps, data engineering, and

related technologies. It combines instructor-led training, hands-on labs, projects, and placement assistance to help learners prepare for real-world roles.

◆ **Highlights**

- Offers **cloud & AI/ML bootcamps**, including AWS, Azure, Google Cloud, and DevOps programs.
- Provides hands-on practice with real projects and labs.
- Career support, resume prep, interview guidance, and community access.

🔗 Official Site: <https://k21academy.com/>

 **Summary Table**

Topic	Definition	Official Link
AWS	Cloud computing platform with over 240 services, scalable, secure, pay-as-you-go.	https://aws.amazon.com/what-is-aws/
AWS Bedrock	Managed generative AI service with access to multiple foundation models.	https://aws.amazon.com/bedrock/
Amazon SageMaker	Managed machine learning platform for building, training, and deploying models.	https://aws.amazon.com/sagemaker/
RAG	Technique combining retrieval from knowledge bases with	https://docs.aws.amazon.com/sagemaker/latest/dg/jumpstart-foundation-models-customize-rag.html

Topic	Definition	Official Link
K21 Academy	generative AI. Tech training platform offering cloud, AI/ML, and DevOps programs.	